

ECE 485/585

Computer Organization and Design

Lecture 3: Instruction Set Architecture

Fall 2022

Won-Jae Yi, Ph.D.

Department of Electrical and Computer Engineering
Illinois Institute of Technology

Instruction Set

- The repertoire of instructions of a computer
- Different computers have different instruction sets
 - But with many aspects in common
- Early computers had very simple instruction sets
 - Simplified implementation
- Many modern computers also have simple instruction sets

The MIPS Instruction Set

- Used as the example throughout the book
- Stanford MIPS commercialized by MIPS Technologies (www.mips.com)
- Large share of embedded core market
 - Applications in consumer electronics, network/storage equipment, cameras, printers, ...
- Typical of many modern ISAs
 - See MIPS Reference Data tear-out card, and Appendixes B and E



Arithmetic Operations operands ..

- Add and subtract, three operands
 - Two sources and one destination

OP ← add a, b, c # a gets b + c

- All arithmetic operations have this form

✓ *Design Principle 1: Simplicity favors regularity*

- Regularity makes implementation simpler (HW)
- Simplicity enables higher performance at lower cost

Arithmetic Example

- C code:

`f = (g + h) - (i + j);`

- Compiled MIPS code:

`add t0, g, h # temp t0 = g + h`
`add t1, i, j # temp t1 = i + j`
`sub f, t0, t1 # f = t0 - t1`

Register Operands

- Arithmetic instructions use register operands
- MIPS has a 32 × 32-bit register file
 - Use for frequently accessed data
 - Numbered 0 to 31
 - 32-bit data called a “word”
 - A fixed-sized data unit by the specific instruction set
- Assembler names
 - \$t0, \$t1, ..., \$t9 for temporary values
 - \$s0, \$s1, ..., \$s7 for saved variables
- *Design Principle 2: Smaller is faster*
 - c.f. main memory: millions of locations

recognized
by
ALU

array

1 word = 32 bits

Register Operand Example

- C code:

f = (g + h) - (i + j);

- f, ..., j in \$s0, ..., \$s4

- Compiled MIPS code:

add \$t0, \$s1, \$s2


add \$t1, \$s3, \$s4

sub \$s0, \$t0, \$t1

Memory Operands

LOAD & STORE Arch.
32 bits

- Main memory used for composite data
 - Arrays, structures, dynamic data
- • To apply arithmetic operations
 - Load values from memory into registers
 - Store result from register to memory
- Memory is byte addressed
 - Each address identifies an 8-bit byte
- • Words are aligned in memory
 - Address must be a multiple of 4 ✓
- MIPS can be either Big Endian or Little Endian
 - Big Endian: Most-significant byte (MSB) at least address of a word
 - Little Endian: Least-significant byte (LSB) at least address of a word



Big Endian

Vs Little Endian

$$0x1000 = 0x\underline{FF}00AA11$$

0x1000	FF	11
0x1001	00	AA
0x1002	AA	00
0x1003	11	FF

8 bits.
B.E.

L.E.

(MIPS) \swarrow B.E.
L.E.

Memory Operand Example 1

- C code: $\$s2$ base addr + offset

$g = h + A[8];$

- g in $\$s1$, h in $\$s2$, base address of A in $\$s3$

- Compiled MIPS code:


- Index 8 requires offset of 32
 - 4 bytes per word

lw $\$t0, 32(\$s3)$ # load word
add $\$s1, \$s2, \$t0$

offset

base register

$0x2020$
2010



$8 \times 4 \rightarrow 2$
 $= 32$ bytes

base addr + 32
16

1000000
 $\Rightarrow 2010$

$\$s3 = 0x2000$

Memory Operand Example 2

- C code:

A[12] = h + A[8];

- h in \$s2, base address of A in \$s3

- Compiled MIPS code:

- Index 8 requires offset of 32

✓ lw \$t0, 32(\$s3) # load word

add \$t0, \$s2, \$t0

✓ sw \$t0, 48(\$s3) # store word

48 = 4 × 12

Register vs. Memory

- Registers are faster to access than memory
- Operating on memory data requires loads and stores
 - More instructions to be executed
- Compiler must use registers for variables as much as possible
 - Only spill to memory for less frequently used variables
 - Register optimization is important!

Immediate Operands

- Constant data specified in an instruction

`addi $s3, $s3, 4`

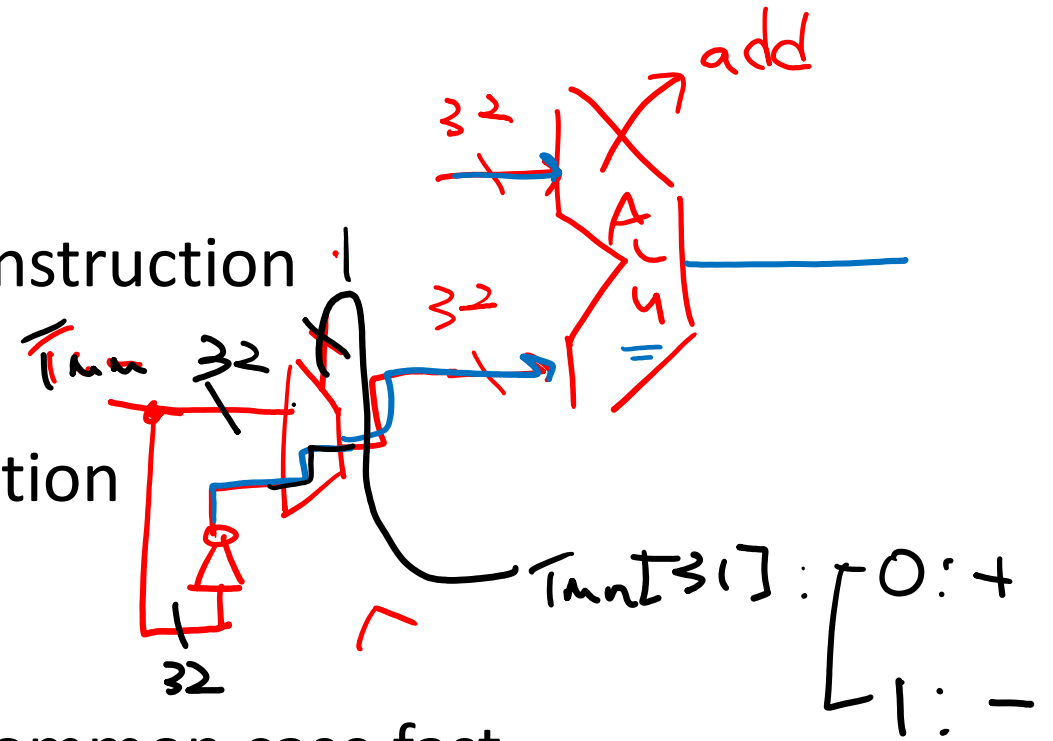
- No subtract immediate instruction

- Just use a negative constant

`addi $s2, $s1, -1`

- Design Principle 3: Make the common case fast

- Small constants are common
- Immediate operand avoids a load instruction



The Constant Zero

- MIPS register 0 (\$zero) is the constant 0
 - Cannot be overwritten
 - Useful for common operations
 - E.g., move between registers
- add \$t2, \$s1, \$zero

~~MOVE~~

Unsigned Binary Integers

- Given an n-bit number

$$x = x_{n-1}2^{n-1} + x_{n-2}2^{n-2} + \dots + x_12^1 + x_02^0$$

- Range: 0 to $+2^n - 1$

Handwritten red note: $0 \sim 2^n - 1$ and a box containing 2^n with a diagonal line through it.

- Example

- $0000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0000\ 1011_2$
 $= 0 + \dots + 1 \times 2^3 + 0 \times 2^2 + 1 \times 2^1 + 1 \times 2^0$
 $= 0 + \dots + 8 + 0 + 2 + 1 = 11_{10}$

- Using 32 bits

- 0 to +4,294,967,295

2^s Complement Signed Integers

- Given an ⁸n-bit number

$$x = -x_{n-1}2^{n-1} + x_{n-2}2^{n-2} + \dots + x_12^1 + x_02^0$$

- Range: -2^{n-1} to $+2^{n-1} - 1$

- Example

- $$\begin{aligned} &1111\ 1111\ 1111\ 1111\ 1111\ 1111\ 1111\ 1100_2 \\ &= -1 \times 2^{31} + 1 \times 2^{30} + \dots + 1 \times 2^2 + 0 \times 2^1 + 0 \times 2^0 \\ &= -2,147,483,648 + 2,147,483,644 = -4_{10} \end{aligned}$$

- Using 32 bits

- $-2,147,483,648$ to $+2,147,483,647$

$-2^7 \sim 2^7 - 1$ $0 \sim 255$
 $-128 \sim 127$

$-2^{n-1} \sim 2^{n-1} - 1$

2^s Complement Signed Integers

- Bit 31 is sign bit
 - 1 for negative numbers
 - 0 for non-negative numbers
- $-(-2^n - 1)$ can't be represented
- Non-negative numbers have the same unsigned and 2s-complement representation
- Some specific numbers
 - • 0: 0000 0000 ... 0000
 - • -1: 1111 1111 ... 1111
 - (• Most-negative: 1000 0000 ... 0000
 - (• Most-positive: 0111 1111 ... 1111

Signed Negation

- Complement and add 1
 - Complement means $1 \rightarrow 0, 0 \rightarrow 1$



$$x + \bar{x} = 1111 \dots 111_2 = -1$$

$$\bar{x} + 1 = -x$$

- Example: negate +2
 - $+2 = 0000 \ 0000 \dots 0010_2$
 - $-2 = 1111 \ 1111 \dots 1101_2 + 1$
 $= 1111 \ 1111 \dots 1110_2$

Sign Extension

- Representing a number using more bits
 - Preserve the numeric value
- In MIPS instruction set
 - addi: extend immediate value
 - lb, lh: extend loaded byte/halfword
 - beq, bne: extend the displacement
- Replicate the sign bit to the left
 - c.f. unsigned values: extend with 0s
- Examples: 8-bit to 16-bit
 - +2: \downarrow 0000 0010 \Rightarrow 0000 0000 0000 0010
 - 2: 1111 1110 \Rightarrow 1111 1111 1111 1110



Extra 16 NOT REQUIRED

Representing Instructions

- Instructions are encoded in binary
 - Called machine code
- MIPS instructions
 - Encoded as 32-bit instruction words
 - Small number of formats encoding operation code (opcode), register numbers, ...
 - Regularity!


→ Register numbers

- \$t0 – \$t7 are reg's 8 – 15
- \$t8 – \$t9 are reg's 24 – 25
- \$s0 – \$s7 are reg's 16 – 23

ASM

Machine.

R-type




32 bits

Arithmetic
op. 3 REGs

Source Source dest

J-type




32 bits

addr \$s0, \$s1, 1
beg \$s0, \$s1, NEXT

displacement

J-type




32 bits.

J LABEL

<<2

MIPS R-Format Instructions



- Instruction fields

- op: operation code (opcode)
- rs: first source register number
- rt: second source register number
- rd: destination register number
- shamt: shift amount (00000 for now)
- funct: function code (extends opcode)

add \$5?, \$5, \$2

dest. source source.

The handwritten example 'add \$5?, \$5, \$2' is annotated with blue arrows. An arrow points from the 'dest.' label to the '\$5?' register, which is circled in red. Another arrow points from the 'source' label to the '\$5' register. A third arrow points from the 'source.' label to the '\$2' register. A large blue arrow also points from the 'rd' field of the instruction format diagram to the '\$5?' register.

MIPS R-Format Example

IR: instruction register

000000 10001 10010 01000 00000 100000

op	rs	rt	rd	shamt	funct
6 bits	5 bits	5 bits	5 bits	5 bits	6 bits

(R) r_d r_s r_t
 → add \$t0, \$s1, \$s2

special	\$s1	\$s2	\$t0	0	add
---------	------	------	------	---	-----

0	17	18	8	0	32
---	----	----	---	---	----

000000	10001	10010	01000	00000	100000
--------	-------	-------	-------	-------	--------

→ 000000100011001001000000100000₂ = 02324020₁₆

MEM

0023240201

PC " Prog. Counter

Hexadecimal


- Base 16
 - Compact representation of bit strings
 - 4 bits per hex digit

0	0000	4	0100	8	1000	C	1100
1	0001	5	0101	9	1001	D	1101
2	0010	6	0110	A	1010	E	1110
3	0011	7	0111	B	1011	F	1111

- Example: ECA8 6420
 - 1110 1100 1010 1000 0110 0100 0010 0000

MIPS I-Format Instructions

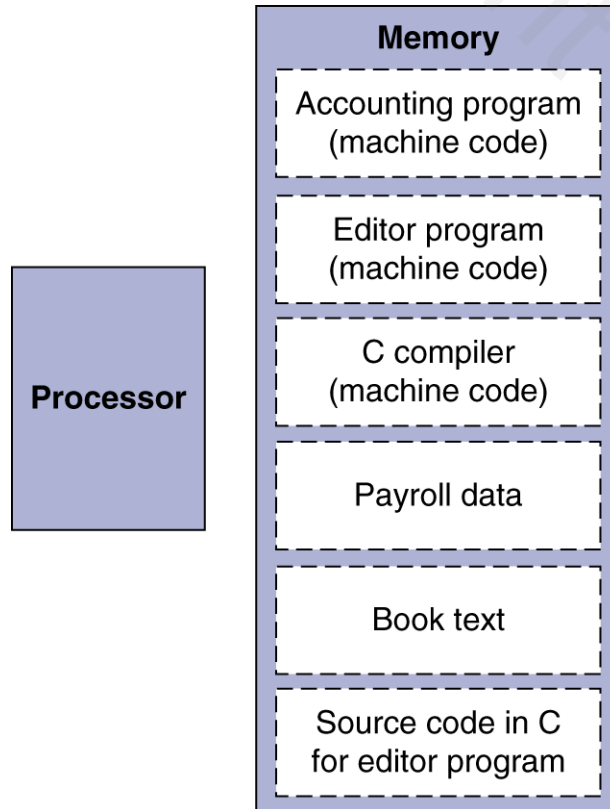
SW \$s0, 32(\$s1)
LW \$s0, 4(\$s1)
↑
rs



- Immediate arithmetic and load/store instructions
 - rt: destination or source register number
 - Constant: -2^{15} to $+2^{15} - 1$
 - Address: offset added to base address in rs
- *Design Principle 4*: Good design demands good compromises
 - Different formats complicate decoding, but allow 32-bit instructions uniformly
 - Keep formats as similar as possible

displacement
- , +

Stored Program Computers



- Instructions represented in binary, just like data
- Instructions and data stored in memory
- Programs can operate on programs
 - e.g., compilers, linkers, ...
- Binary compatibility allows compiled programs to work on different computers
 - Standardized ISAs

Logical Operations

- Instructions for bitwise manipulation

Operation	C	Java	MIPS
Shift left	<<	<<	<u>sll</u>
Shift right	>>	>>>	<u>srl</u>
Bitwise AND	&	&	and, andi
Bitwise OR			or, ori
Bitwise NOT	~	~	nor

- Useful for extracting and inserting groups of bits in a word

Shift Operations



op	rs	rt	rd	shamt	funct
6 bits	5 bits	5 bits	5 bits	5 bits	6 bits

- shamt: how many positions to shift
- Shift left logical
 - Shift left and fill with 0 bits
 - sll by i bits multiplies by 2^i
- Shift right logical
 - Shift right and fill with 0 bits
 - srl by i bits divides by 2^i (unsigned only)

AND Operations

- Useful to mask bits in a word
 - Select some bits, clear others to 0
 - Check on the status of a bit (1) or to force mask off (0)

and \$t0, ~~\$t1~~, \$t2


\$t2	0000 0000 0000 0000 0000 1101 1100 0000
\$t1	0000 0000 0000 0000 0011 1100 0000 0000
\$t0	0000 0000 0000 0000 0000 1100 0000 0000

MASK OFF

\$t1 : 0x00003C00
\$t2 : 0x0000DC00



OR Operations



- Useful to include bits in a word
 - Set some bits to 1, leave others unchanged
 - To leave a bit unchanged (0) or to turn bits ON (1)

or \$t0, \$t1, \$t2

→ MASK ON

\$t2	0000 0000 0000 0000 0000 1101 1100 0000
------	---

\$t1	0000 0000 0000 0000 0011 1100 0000 0000
------	---

\$t0	0000 0000 0000 0000 0011 1101 1100 0000
------	---

NOT Operations

- Useful to invert bits in a word
 - Change 0 to 1, and 1 to 0
- MIPS has NOR 3-operand instruction
 - $a \text{ NOR } b == \text{NOT} (a \text{ OR } b)$

→ `nor $t0, $t1, $zero`

Register 0:
always read as
zero

`$t1` 0000 0000 0000 0000 0011 1100 0000 0000

`$t0` 1111 1111 1111 1111 1100 0011 1111 1111

Conditional Operations

- Branch to a labeled instruction if a condition is true
 - Otherwise, continue sequentially
- beq (rs, rt), L1
 - if ($rs == rt$) branch to instruction labeled L1;
- bne (rs), (rt), L1
 - if ($rs != rt$) branch to instruction labeled L1;
- j L1
 - unconditional jump to instruction labeled L1

Extra
1/2

Compiling If Statements

- C code:


\$s3 \$s4 \$s0 \$s1 \$s2
if (i==j) f = g+h;
else f = g-h;

- f, g, ... in \$s0, \$s1, ...

- Compiled MIPS code:

bne \$s3, \$s4, Else
add \$s0, \$s1, \$s2
j Exit
Else: sub \$s0, \$s1, \$s2
Exit: ...

Assembler calculates addresses



bne
beg

Compiling Loop Statements

- C code:


while (save[i] == k) i += 1;

- i in \$s3, k in \$s5, address of save in \$s6

- Compiled MIPS code:

```

Loop: sll $t1, $s3, 2
      add $t1, $t1, $s6
      lw  $t0, 0($t1)
      bne $t0, $s5, Exit
      addi $s3, $s3, 1
      j   Loop
Exit: ...
    
```




Basic Blocks

- A basic block is a sequence of instructions with
 - No embedded branches (except at end)
 - No branch targets (except at beginning)

array[8]

→ array[0]
c2]



Python

- A compiler identifies basic blocks for optimization
- An advanced processor can accelerate execution of basic blocks

More Conditional Operations

- Set result to 1 if a condition is true
 - Otherwise, set to 0
- `slt rd, rs, rt`
 - if ($rs < rt$) $rd = 1$; else $rd = 0$;
- `slti rt, rs, constant`
 - if ($rs < \text{constant}$) $rt = 1$; else $rt = 0$;
- Use in combination with `beq`, `bne`
`slt $t0, $s1, $s2 # if ($s1 < $s2)`
`bne $t0, $zero, L # branch to L`

Branch Instruction Design

- Why not **blt**, **bge**, etc?
- Hardware for $<$, \geq , ... slower than $=$, \neq
 - Combining with branch involves more work per instruction, requiring a slower clock
 - All instructions penalized!
- **beq** and **bne** are the common case
- This is a good design compromise

Signed vs. Unsigned

- Signed comparison: `slt`, `slti`
- Unsigned comparison: `sltu`, `sltui`
- Example
 - `$s0 = 1111 1111 1111 1111 1111 1111 1111 1111`
 - `$s1 = 0000 0000 0000 0000 0000 0000 0000 0001`
 - `slt $t0, $s0, $s1 # signed`
 - $-1 < +1 \Rightarrow \$t0 = 1$
 - `sltu $t0, $s0, $s1 # unsigned`
 - $+4,294,967,295 > +1 \Rightarrow \$t0 = 0$